


Gaussian processes for time series

- Anna Scaife's (University of Manchester) blog:
<https://allofyourbases.com/2017/08/21/gaussian-processes-in-python/>
<https://allofyourbases.com/2017/09/04/gaussian-process-modelling-in-python/>
<https://allofyourbases.com/2017/09/17/predicting-the-future/>
<https://allofyourbases.com/2017/09/23/spinning-stars/>
<https://allofyourbases.com/2017/10/02/spinning-stars-ii-celerite/>
- Roberts S, Osborne M, Ebdon M, Reece S, Gibson N, Aigrain S. 2013 Gaussian processes for time-series modelling. Phil Trans R Soc A 371: 20110550. <http://dx.doi.org/10.1098/rsta.2011.0550>
- <http://www.gaussianprocess.org/>

Co-variate Gaussian noise

covariance matrix

$$\mathbf{K}(\mathbf{x}, \mathbf{x}) = \begin{pmatrix} k(x_1, x_1) & k(x_1, x_2) & \cdots & k(x_1, x_n) \\ k(x_2, x_1) & k(x_2, x_2) & \cdots & k(x_2, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ k(x_n, x_1) & k(x_n, x_2) & \cdots & k(x_n, x_n) \end{pmatrix}.$$



Co-variate Gaussian noise

$$\mathbf{K}(\mathbf{x}, \mathbf{x}) = \begin{pmatrix} k(x_1, x_1) & 0 & \dots & 0 \\ 0 & k(x_2, x_2) & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & k(x_n, x_n) \end{pmatrix}$$

diagonal co-variance matrix shows
uncorrelated (independent) Gaussian noise

each value in this covariance matrix
represents the variance at some point.

The actual y-value at each x-position is drawn from a probability distribution:

$$p(y_i) = N(\mu_i, K_{ij})$$

If the co-variance matrix is diagonal this becomes:

$$y_i = \mu_i + N(0, \sigma_i^2)$$

But, if the co-variance matrix is non-diagonal, the y-values at one position affects those at another.

The co-variance kernel

- This is the “formula” for making the co-variance matrix
- It controls how the data points affect each other
- Squared-exponential co-variance kernel:

$$k(x_n, x_m) = h^2 \exp \left(\frac{-(x_n - x_m)^2}{\lambda^2} \right)$$

- h (normalisation) and λ (width) are hyper parameters
- There are many other options for co-variance kernels, even periodic ones.

As part of Tut 4:

1. Use a squared-exponential kernel to make a co-variance matrix (300 x 300)
2. Plot the covariance matrix for $\lambda = 0.1, 1$ and 10, with a constant normalisation.
3. Plot the y values (just as a function of index) drawn from this distribution.
4. Comment on the difference between the distributions for different values of λ .
5. Extend this to a model, by splitting your data into test and training sets...

Modelling with Gaussian processes

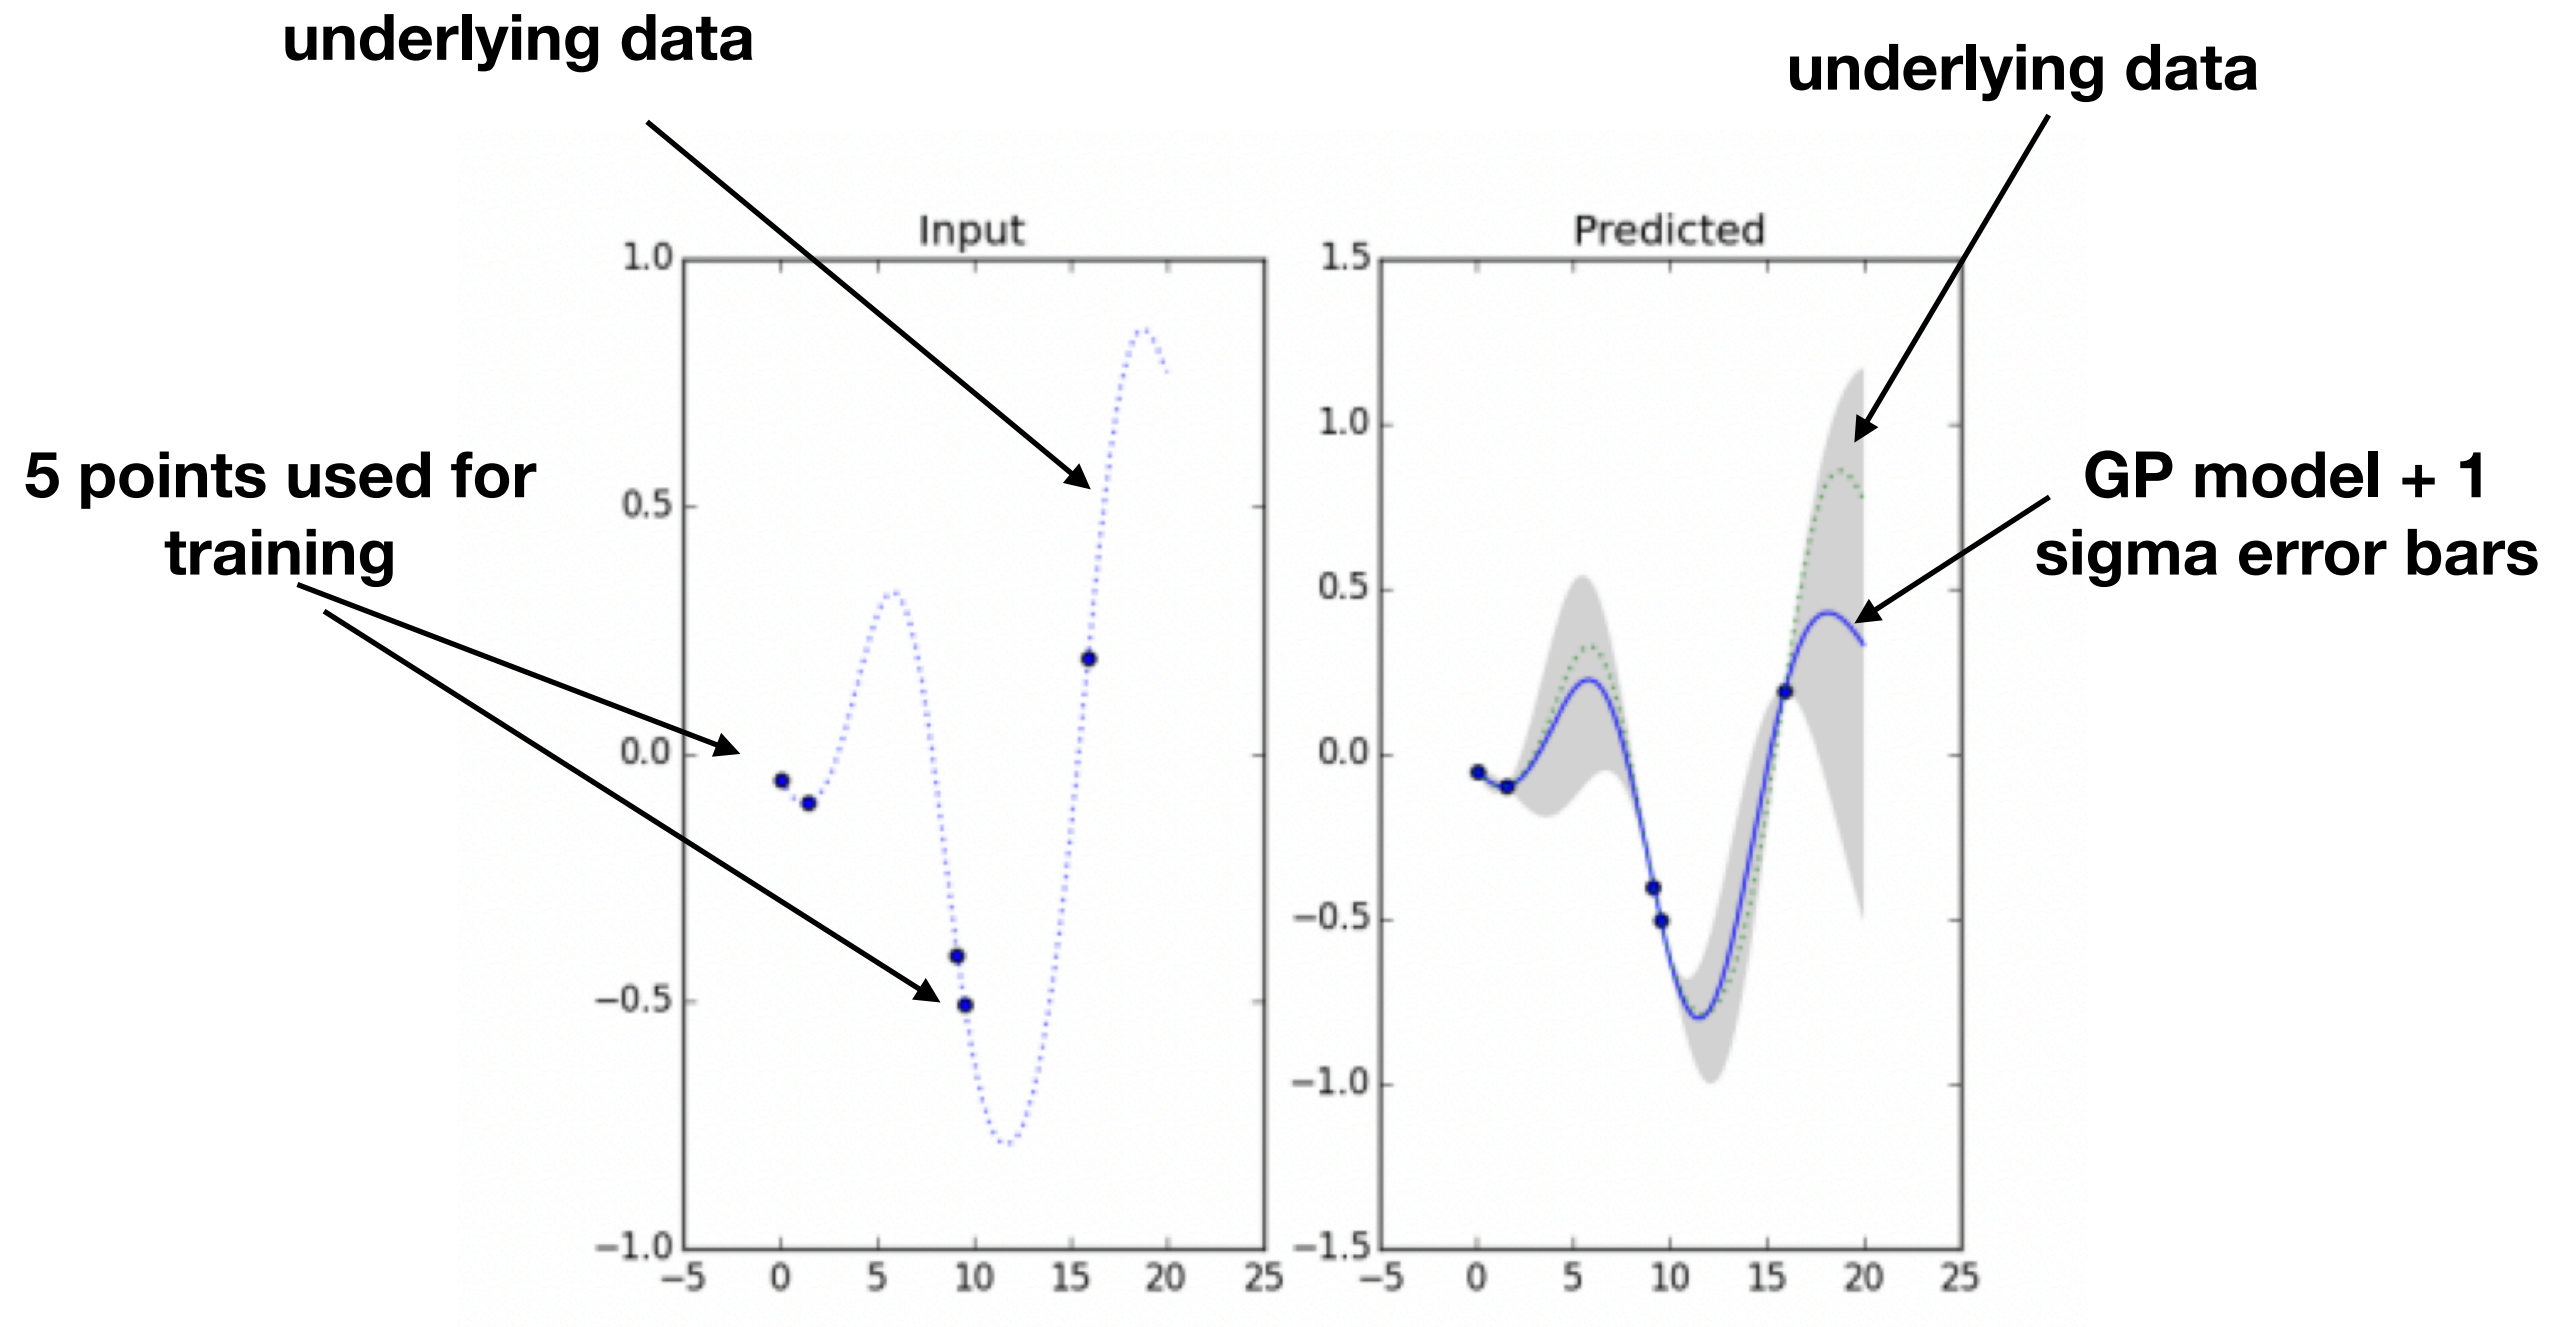
- Let's try to predict the value of a function at any point based on its value at a few known points
- Assumption: *behaviour of function depends only on the covariance between measurement at different separations*
- Invert the probability statement to get y_i at some x_i : $p(y_i) = N(\mu_i, K_{ij})$
- You need to know or calculate covariance matrix and some training data.

$$\begin{aligned} \mathbf{m}_* &= \mathbf{K}(\mathbf{x}_*, \mathbf{x})^T \mathbf{K}(\mathbf{x}, \mathbf{x})^{-1} \mathbf{y} \\ \mathbf{C}_* &= \mathbf{K}(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{K}(\mathbf{x}_*, \mathbf{x})^T \mathbf{K}(\mathbf{x}, \mathbf{x})^{-1} \mathbf{K}(\mathbf{x}, \mathbf{x}_*) \end{aligned}$$

This is for measurements with zero mean

$$\begin{aligned} \mathbf{m}_* &= \mathbf{k}_*^T \mathbf{K}^{-1} \mathbf{y} \\ \mathbf{C}_* &= \mathbf{k}(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{k}_*^T \mathbf{K}^{-1} \mathbf{k}_* \end{aligned}$$

Modelling with Gaussian processes



Markov chain Monte Carlo

- Useful method for integrating over a large number of parameters
- Monte Carlo - statistically independent samples
- MCMC samples from continuous random variable, with prob density proportional to a known function
- Generate an ensemble of chains starting from arbitrarily chosen points. Move about randomly according to an algorithm that searches for where high contributions to the integral are.
- Metropolis-Hastings algorithm is a well-known implementation of the random walker.